

# Decision Making and Reinforcement Learning

## Module 7: Temporal-Difference Learning

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences

# Topics

---

- Temporal difference prediction
- Batch learning methods
- SARSA (on-policy control)
- Q-learning (off-policy control)

# Learning Objectives

---

- **Implement** TD prediction for a given policy
- **Apply** TD methods in batch and/or offline
- **Implement** the SARSA and Q-learning algorithms
- **Compare and contrast** properties and characteristics of on-policy vs off-policy TD learning

# Review: Monte Carlo Prediction

---

- *Prediction*: Estimate state values for given policy  $\pi$
- MC methods average observed utilities from multiple episodes

$$V^{\pi}(s) = \frac{1}{N} \sum_{i=1}^N G_i(s)$$

# Review: Monte Carlo Prediction

---

- *Prediction*: Estimate state values for given policy  $\pi$
- MC methods average observed utilities from multiple episodes

$$V^{\pi}(s) = \frac{1}{N} \sum_{i=1}^N G_i(s)$$

- In our implementation, we used a “moving average” calculation
- Suppose  $N(s)$  is the number of visits to state  $s$  prior to current episode

$$V^{\pi}(s) \leftarrow \frac{N(s)V^{\pi}(s) + G_i(s)}{N(s) + 1}$$

- Let's rewrite the state value update expression...

$$V^\pi(s) \leftarrow \frac{N(s)V^\pi(s) + G_i(s)}{N(s) + 1} = V^\pi(s) + \frac{1}{N(s) + 1} (G_i(s) - V^\pi(s))$$

- “Old estimate” + “step size”  $\times$  (“target” – “old estimate”)

- Let's rewrite the state value update expression...

$$V^\pi(s) \leftarrow \frac{N(s)V^\pi(s) + G_i(s)}{N(s) + 1} = V^\pi(s) + \frac{1}{N(s) + 1} (G_i(s) - V^\pi(s))$$

- “Old estimate” + “step size”  $\times$  (“target” – “old estimate”)
- We can also choose to give more weight to *recent* returns
- **Constant- $\alpha$  MC** *exponentially* decays the weights on past returns by learning rate  $\alpha$

$$V^\pi(s) \leftarrow V^\pi(s) + \alpha (G_i(s) - V^\pi(s))$$

# Example: Constant- $\alpha$ MC

---

- Suppose  $V(s) = 0$ ; consider the reward sequence  $+3, -2, +1$
- Moving average MC:
- $V(s) = 0 + \frac{1}{1}(3 - 0) = 3$
- $V(s) = 3 + \frac{1}{2}(-2 - 3) = \frac{1}{2}$
- $V(s) = \frac{1}{2} + \frac{1}{3}\left(1 - \frac{1}{2}\right) = \frac{2}{3}$



# Example: Constant- $\alpha$ MC

- Suppose  $V(s) = 0$ ; consider the reward sequence  $+3, -2, +1$
- Moving average MC:
  - $V(s) = 0 + \frac{1}{1}(3 - 0) = 3$
  - $V(s) = 3 + \frac{1}{2}(-2 - 3) = \frac{1}{2}$
  - $V(s) = \frac{1}{2} + \frac{1}{3}\left(1 - \frac{1}{2}\right) = \frac{2}{3}$
- Constant  $\alpha$  MC,  $\alpha = 0.5$ :
  - $V(s) = 0 + \frac{1}{2}(3 - 0) = 1.5$
  - $V(s) = 1.5 + \frac{1}{2}(-2 - 1.5) = -0.25$
  - $V(s) = -0.25 + \frac{1}{2}(1 + 0.25) = 0.375$

# Temporal-Difference Learning

---

- MC requires episodic structure—what about infinite horizon problems?
- Recall from DP:  $V^\pi(s)$  depends on values of successors of  $s$

# Temporal-Difference Learning

- MC requires episodic structure—what about infinite horizon problems?
- Recall from DP:  $V^\pi(s)$  depends on values of successors of  $s$
- **One-step TD ( $TD(0)$ ):** Replace the *target* term with the sum of immediate reward with discounted successor state value!


$$V^\pi(s) \leftarrow V^\pi(s) + \alpha \left( \underbrace{r_{t+1} + \gamma V^\pi(s')}_{\text{Target}} - V^\pi(s) \right)$$

TD error  $\delta_t$

- 
- **Given:** Policy  $\pi$ , *learning rate*  $\alpha$  between 0 and 1
  - **Initialize**  $V^\pi(s) \leftarrow 0$

- 
- **Given:** Policy  $\pi$ , *learning rate*  $\alpha$  between 0 and 1
  - **Initialize**  $V^\pi(s) \leftarrow 0$
  - **Loop:**
    - **Initialize** starting state  $s$  if needed
    - **Generate** sequence  $(s, \pi(s), r, s')$





# Example: Mini-Gridworld

---

- All values initialized to 0;  $\gamma = 0.8$ ,  $\alpha = 0.5$
- Policy to evaluate:  $\pi(s) = L$  for all states
- Observed state and reward sequence:  $(A, +3, A, -2, B, +1, C, -2, B, +3, A)$





# Example: Mini-Gridworld

- All values initialized to 0;  $\gamma = 0.8$ ,  $\alpha = 0.5$
- Policy to evaluate:  $\pi(s) = L$  for all states



- Observed state and reward sequence:  $(A, +3, A, -2, B, +1, C, -2, B, +3, A)$

Transition	
	0
	0

$$V^{\pi}(A) \leftarrow V^{\pi}(A) + \alpha(r + \gamma V^{\pi}(A) - V^{\pi}(A))$$

$$V^{\pi}(A) \leftarrow 0 + 0.5(3 + 0.8(0) - 0)$$

# Example: Mini-Gridworld

- All values initialized to 0;  $\gamma = 0.8$ ,  $\alpha = 0.5$
- Policy to evaluate:  $\pi(s) = L$  for all states
- Observed state and reward sequence:  $(A, +3, A, -2, B, +1, C, -2, B, +3, A)$



Transition		
	0	0
	0	0

$$V^\pi(A) \leftarrow V^\pi(A) + \alpha(r + \gamma V^\pi(B) - V^\pi(A))$$

$$V^\pi(A) \leftarrow 1.5 + 0.5(-2 + 0.8(0) - 1.5)$$

# Example: Mini-Gridworld

- All values initialized to 0;  $\gamma = 0.8$ ,  $\alpha = 0.5$
- Policy to evaluate:  $\pi(s) = L$  for all states



- Observed state and reward sequence:  $(A, +3, A, -2, B, +1, C, -2, B, +3, A)$

Transition			
	0	0	
	0	0	0

$$V^\pi(B) \leftarrow 0 + 0.5(1 + 0.8(0) - 0)$$

# Example: Mini-Gridworld

- All values initialized to 0;  $\gamma = 0.8$ ,  $\alpha = 0.5$
- Policy to evaluate:  $\pi(s) = L$  for all states
- Observed state and reward sequence:  $(A, +3, A, -2, B, +1, C, -2, B, +3, A)$



Transition					
	0	0			
	0	0	0		

- 
- TD methods perform updates immediately with no episodic structure (MC)
  - Useful if problems have long episodes or are continuing tasks



# Offline Prediction

---

- We've presented constant- $\alpha$  MC and  $TD(0)$  as online, sample-based algorithms in contrast to dynamic programming
- It is also possible to use these as *offline*, sample-based algorithms

# Offline Prediction

---

- We've presented constant- $\alpha$  MC and  $TD(0)$  as online, sample-based algorithms in contrast to dynamic programming
- It is also possible to use these as *offline*, sample-based algorithms
- Suppose we have a fixed set of samples in place of underlying model
- We can run  $TD(0)$  on the samples over and over until values converge



# Offline Prediction

---

- We've presented constant- $\alpha$  MC and  $TD(0)$  as online, sample-based algorithms in contrast to dynamic programming
- It is also possible to use these as *offline*, sample-based algorithms
- Suppose we have a fixed set of samples in place of underlying model
- We can run  $TD(0)$  on the samples over and over until values converge
- We can perform **batch updating**, in which we make a single TD update based on the *sum* of all sample TD errors in sample set

# Batch Updating

---

- **Given:** Training data  $s_0, a_0, r_1, \dots, s_{T-1}, a_{T-1}, r_T$ ; learning rate  $\alpha$
- **Initialize**  $V^\pi(s) \leftarrow 0$

# Batch Updating

---

- **Given:** Training data  $s_0, a_0, r_1, \dots, s_{T-1}, a_{T-1}, r_T$ ; learning rate  $\alpha$
- **Initialize**  $V^\pi(s) \leftarrow 0$
- **Loop** until values converge:
  - $\delta(s) \leftarrow 0 \ \forall s$
  - **For** each state  $s_t$  in training data:
    - $\delta(s_t) \leftarrow \delta(s_t) + (r_{t+1} + \gamma V(s_{t+1}) - V(s_t))$

# Batch Updating

- **Given:** Training data  $s_0, a_0, r_1, \dots, s_{T-1}, a_{T-1}, r_T$ ; learning rate  $\alpha$
- **Initialize**  $V^\pi(s) \leftarrow 0$
- **Loop** until values converge:
  - $\delta(s) \leftarrow 0 \ \forall s$
  - **For** each state  $s_t$  in training data:
    - $\delta(s_t) \leftarrow \delta(s_t) + (r_{t+1} + \gamma V(s_{t+1}) - V(s_t))$
  - **For** each state  $s$ :
    - $V^\pi(s) \leftarrow V^\pi(s) + \alpha \delta(s)$

# Example: Mini-Gridworld

- All values initialized to 0;  $\gamma = 0.8$ ,  $\alpha = 0.5$
- Policy to evaluate:  $\pi(s) = L$  for all states
- Observed state and reward sequence:  $(A, +3, A, -2, B, +1, C, -2, B, +3, A)$



$$\delta(s_t) \leftarrow r_{t+1} + \gamma V^\pi(s_{t+1}) - V^\pi(s_t)$$

$$V^\pi(s_t) \leftarrow V^\pi(s_t) + \alpha \delta(s_t)$$

# Example: Mini-Gridworld

- All values initialized to 0;  $\gamma = 0.8$ ,  $\alpha = 0.5$
- Policy to evaluate:  $\pi(s) = L$  for all states
- Observed state and reward sequence:  $(A, +3, A, -2, B, +1, C, -2, B, +3, A)$



$$\delta(s_t) \leftarrow r_{t+1} + \gamma V^\pi(s_{t+1}) - V^\pi(s_t)$$

$$V^\pi(s_t) \leftarrow V^\pi(s_t) + \alpha \delta(s_t)$$

Iteration			

$$\begin{aligned} \delta(A) \leftarrow & (3 + \gamma V(A) - V(A)) \\ & + (-2 + \gamma V(B) - V(A)) \end{aligned}$$

$$V(A) \leftarrow V(A) + \alpha \delta(A)$$

# Example: Mini-Gridworld

- All values initialized to 0;  $\gamma = 0.8$ ,  $\alpha = 0.5$
- Policy to evaluate:  $\pi(s) = L$  for all states
- Observed state and reward sequence:  $(A, +3, A, -2, B, +1, C, -2, B, +3, A)$



$$\delta(s_t) \leftarrow r_{t+1} + \gamma V^\pi(s_{t+1}) - V^\pi(s_t)$$

$$V^\pi(s_t) \leftarrow V^\pi(s_t) + \alpha \delta(s_t)$$

Iteration			

$$\begin{aligned} \delta(B) \leftarrow & (1 + \gamma V(C) - V(B)) \\ & + (3 + \gamma V(A) - V(B)) \end{aligned}$$

$$V(B) \leftarrow V(B) + \alpha \delta(B)$$

# Example: Mini-Gridworld

- All values initialized to 0;  $\gamma = 0.8$ ,  $\alpha = 0.5$
- Policy to evaluate:  $\pi(s) = L$  for all states



- Observed state and reward sequence:  $(A, +3, A, -2, B, +1, C, -2, B, +3, A)$

$$\delta(s_t) \leftarrow r_{t+1} + \gamma V^\pi(s_{t+1}) - V^\pi(s_t)$$

$$V^\pi(s_t) \leftarrow V^\pi(s_t) + \alpha \delta(s_t)$$

Iteration			

$$\delta(C) \leftarrow -2 + \gamma V(B) - V(C)$$

$$V(C) \leftarrow V(C) + \alpha \delta(C)$$



# Example: Mini-Gridworld

- All values initialized to 0;  $\gamma = 0.8$ ,  $\alpha = 0.5$
- Policy to evaluate:  $\pi(s) = L$  for all states



- Observed state and reward sequence:  $(A, +3, A, -2, B, +1, C, -2, B, +3, A)$

$$\delta(s_t) \leftarrow r_{t+1} + \gamma V^\pi(s_{t+1}) - V^\pi(s_t)$$

$$V^\pi(s_t) \leftarrow V^\pi(s_t) + \alpha \delta(s_t)$$

Iteration					

# Example: Mini-Gridworld

- All values initialized to 0;  $\gamma = 0.8$ ,  $\alpha = 0.5$
- Policy to evaluate:  $\pi(s) = L$  for all states
- Observed state and reward sequence:  $(A, +3, A, -2, B, +1, C, -2, B, +3, A)$

$$\delta(s_t) \leftarrow r_{t+1} + \gamma V^\pi(s_{t+1}) - V^\pi(s_t)$$

$$V^\pi(s_t) \leftarrow V^\pi(s_t) + \alpha \delta(s_t)$$

[illegible]

# TD Learning for Control

---

- We can apply TD updates to Q values  $Q(s, a)$  rather than state values  $V(s)$
- After each transition  $(s, a, r)$ , we apply the TD update to  $Q(s, a)$

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma Q(s', a') - Q(s, a))$$

# TD Learning for Control

---

- We can apply TD updates to Q values  $Q(s, a)$  rather than state values  $V(s)$
- After each transition  $(s, a, r)$ , we apply the TD update to  $Q(s, a)$

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma Q(s', a') - Q(s, a))$$

- Action  $a$  in state  $s$  has been chosen according to behavior policy
- May allow for exploration to learn about different actions

# TD Learning for Control

---

- We can apply TD updates to Q values  $Q(s, a)$  rather than state values  $V(s)$
- After each transition  $(s, a, r)$ , we apply the TD update to  $Q(s, a)$

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma Q(s', a') - Q(s, a))$$

- Action  $a$  in state  $s$  has been chosen according to behavior policy
- May allow for exploration to learn about different actions
- Which successor state Q value  $Q(s', a')$  to use?
- As with MC control, we need to consider *on-policy* vs *off-policy* learning

# On-Policy: SARSA

---

- **Given:** Step size  $\alpha$ , exploration rate  $\varepsilon$
- **Initialize**  $Q(s, a) \leftarrow 0$ , behavior policy  $\pi$  (e.g.,  $\varepsilon$ -greedy)

# On-Policy: SARSA

---

- **Given:** Step size  $\alpha$ , exploration rate  $\varepsilon$
- **Initialize**  $Q(s, a) \leftarrow 0$ , behavior policy  $\pi$  (e.g.,  $\varepsilon$ -greedy)
- **Loop:**
  - **Initialize** starting state  $s$ , action  $a = \pi(s)$  *if needed*

# On-Policy: SARSA


---

- **Given:** Step size  $\alpha$ , exploration rate  $\varepsilon$
- **Initialize**  $Q(s, a) \leftarrow 0$ , behavior policy  $\pi$  (e.g.,  $\varepsilon$ -greedy)
- **Loop:**
  - **Initialize** starting state  $s$ , action  $a = \pi(s)$  *if needed*
  - **Generate** sequence  $(s, a, r, s')$ ,  $a' \leftarrow \pi(s')$




# On-Policy: SARSA

- **Given:** Step size  $\alpha$ , exploration rate  $\epsilon$
- **Initialize**  $Q(s, a) \leftarrow 0$ , behavior policy  $\pi$  (e.g.,  $\epsilon$ -greedy)
- **Loop:**
  - **Initialize** starting state  $s$ , action  $a = \pi(s)$  *if needed*
  - **Generate** sequence  $(s, a, r, s')$ ,  $a' \leftarrow \pi(s')$
  - $Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma Q(s', a') - Q(s, a))$



# On-Policy: SARSA

- **Given:** Step size  $\alpha$ , exploration rate  $\varepsilon$
- **Initialize**  $Q(s, a) \leftarrow 0$ , behavior policy  $\pi$  (e.g.,  $\varepsilon$ -greedy)
- **Loop:**
  - **Initialize** starting state  $s$ , action  $a = \pi(s)$  *if needed*
  - **Generate** sequence  $(s, a, r, s')$ ,  $a' \leftarrow \pi(s')$
  - $Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma Q(s', a') - Q(s, a))$
  - $s \leftarrow s', a \leftarrow a'$



# Off-Policy: Q-Learning

---

- **Given:** Step size  $\alpha$ , exploration rate  $\varepsilon$
- **Initialize**  $Q(s, a) \leftarrow 0$ , behavior policy  $\pi$  (e.g.,  $\varepsilon$ -greedy)


# Off-Policy: Q-Learning

---

- **Given:** Step size  $\alpha$ , exploration rate  $\varepsilon$
- **Initialize**  $Q(s, a) \leftarrow 0$ , behavior policy  $\pi$  (e.g.,  $\varepsilon$ -greedy)
- **Loop:**
  - **Initialize** starting state  $s$  *if needed*, action  $a = \pi(s)$
  - **Generate** sequence  $(s, a, r, s')$


# Off-Policy: Q-Learning

- **Given:** Step size  $\alpha$ , exploration rate  $\varepsilon$
- **Initialize**  $Q(s, a) \leftarrow 0$ , behavior policy  $\pi$  (e.g.,  $\varepsilon$ -greedy)
- **Loop:**
  - **Initialize** starting state  $s$  *if needed*, action  $a = \pi(s)$
  - **Generate** sequence  $(s, a, r, s')$
  - $Q(s, a) \leftarrow Q(s, a) + \alpha \left( r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right)$



# Off-Policy: Q-Learning

- **Given:** Step size  $\alpha$ , exploration rate  $\varepsilon$
- **Initialize**  $Q(s, a) \leftarrow 0$ , behavior policy  $\pi$  (e.g.,  $\varepsilon$ -greedy)
- **Loop:**
  - **Initialize** starting state  $s$  *if needed*, action  $a = \pi(s)$
  - **Generate** sequence  $(s, a, r, s')$
  - $Q(s, a) \leftarrow Q(s, a) + \alpha \left( r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right)$
  - $s \leftarrow s'$



# Example: Mini-Gridworld

---


- Suppose we currently have  $Q(A, L) = 1.5$ ,  $Q(A, R) = 1$
- Behavior policy is  $\epsilon$ -greedy;  $\alpha = 0.5$ ,  $\gamma = 0.8$



# Example: Mini-Gridworld

---


- Suppose we currently have  $Q(A, L) = 1.5$ ,  $Q(A, R) = 1$
- Behavior policy is  $\epsilon$ -greedy;  $\alpha = 0.5$ ,  $\gamma = 0.8$
- Observed  $(s, a, r, s')$  sequence:  $A, L, +3, A$
- Suppose behavior policy generates  $a' = R$  (*explore*)






# Example: Mini-Gridworld

- Suppose we currently have  $Q(A, L) = 1.5$ ,  $Q(A, R) = 1$
- Behavior policy is  $\epsilon$ -greedy;  $\alpha = 0.5$ ,  $\gamma = 0.8$
- Observed  $(s, a, r, s')$  sequence:  $A, L, +3, A$
- Suppose behavior policy generates  $a' = R$  (*explore*)
- SARSA target:  $r + \gamma Q(A, R) = 3 + 0.8(1) = 3.8$
- New Q-value:  $Q(A, L) \leftarrow 1 + 0.5(3.8 - 1) = 2.4$



# Example: Mini-Gridworld

- Suppose we currently have  $Q(A, L) = 1.5$ ,  $Q(A, R) = 1$
- Behavior policy is  $\epsilon$ -greedy;  $\alpha = 0.5$ ,  $\gamma = 0.8$
- Observed  $(s, a, r, s')$  sequence:  $A, L, +3, A$
- Suppose behavior policy generates  $a' = R$  (*explore*)
- Q-learning target:  $r + \gamma \max_a Q(A, a) = r + \gamma Q(A, L) = 3 + 0.8(1.5) = 4.2$
- New Q-value:  $Q(A, L) \leftarrow 1 + 0.5(4.2 - 1) = 2.6$



# SARSA vs Q-Learning

SARSA	Q-learning
<ul style="list-style-type: none"><li>• On-policy: learn values of behavior policy</li></ul>	<ul style="list-style-type: none"><li>• Off-policy: learn values of target policy</li></ul>
<ul style="list-style-type: none"><li>•</li></ul>	<ul style="list-style-type: none"><li>•</li></ul>
<ul style="list-style-type: none"><li>•</li></ul>	

# SARSA vs Q-Learning


SARSA	Q-learning
<ul style="list-style-type: none"><li>• On-policy: learn values of behavior policy</li></ul>	<ul style="list-style-type: none"><li>• Off-policy: learn values of target policy</li></ul>
<ul style="list-style-type: none"><li>• Agents are more “cautious” as they worry about low rewards from exploration</li></ul>	<ul style="list-style-type: none"><li>• Agents are more “optimistic” as only the best (greedy) actions at each state matter</li></ul>
<ul style="list-style-type: none"><li>• </li></ul>	

# SARSA vs Q-Learning

SARSA	Q-learning
<ul style="list-style-type: none"><li>• On-policy: learn values of behavior policy</li></ul>	<ul style="list-style-type: none"><li>• Off-policy: learn values of target policy</li></ul>
<ul style="list-style-type: none"><li>• Agents are more “cautious” as they worry about low rewards from exploration</li></ul>	<ul style="list-style-type: none"><li>• Agents are more “optimistic” as only the best (greedy) actions at each state matter</li></ul>
<ul style="list-style-type: none"><li>• SARSA and Q-learning are identical if behavior and target policy are the same, or if exploration is removed!</li></ul>	


# Cliff Walking (Sutton and Barto, Ex. 6.6)

- Deterministic gridworld with one terminal (goal) state
- Living reward is  $-1$  in all states except for “cliff”, which rewards  $-100$




# Cliff Walking (Sutton and Barto, Ex. 6.6)

- Deterministic gridworld with one terminal (goal) state
- Living reward is  $-1$  in all states except for “cliff”, which rewards  $-100$
- Since transitions are deterministic, the optimal action at each state is to head in goal direction while ignoring the cliff
- We can use TD control to learn Q values and the corresponding policy



# Cliff Walking (Sutton and Barto, Ex. 6.6)


- Suppose we use SARSA (on-policy) with exploratory behavior policy
- States nearer cliff will have lower Q values than those further away!






# Cliff Walking (Sutton and Barto, Ex. 6.6)

- Suppose we use SARSA (on-policy) with exploratory behavior policy
- States nearer cliff will have lower Q values than those further away!
- Exploration introduces stochasticity into the problem
- Agent learns that it is more likely to “fall” into the cliff if closer to it




# Cliff Walking (Sutton and Barto, Ex. 6.6)

- Suppose we use SARSA (on-policy) with exploratory behavior policy
- States nearer cliff will have lower Q values than those further away!
- Exploration introduces stochasticity into the problem
- Agent learns that it is more likely to “fall” into the cliff if closer to it
- SARSA ends up learning the “safer path”
- Get to the goal but minimize risk of falling




# Cliff Walking (Sutton and Barto, Ex. 6.6)

- Suppose we use Q-learning (off-policy)
- Converged Q-values will not depend on proximity to cliff




# Cliff Walking (Sutton and Barto, Ex. 6.6)

- Suppose we use Q-learning (off-policy)
- Converged Q-values will not depend on proximity to cliff
- Even if behavior policy is exploratory, the Q-value update is “optimistic” and assumes that the agent acts greedily at all times




# Cliff Walking (Sutton and Barto, Ex. 6.6)

- Suppose we use Q-learning (off-policy)
- Converged Q-values will not depend on proximity to cliff
- Even if behavior policy is exploratory, the Q-value update is “optimistic” and assumes that the agent acts greedily at all times
- Learned Q-values will be optimal
- Shortest path will result from learned policy



# MDP Method Comparison


<https://www.davidsilver.uk/wp-content/uploads/2020/03/MC-TD.pdf>



Dynamic programming


# MDP Method Comparison

<https://www.davidsilver.uk/wp-content/uploads/2020/03/MC-TD.pdf>




Dynamic programming

Monte Carlo




# MDP Method Comparison

<https://www.davidsilver.uk/wp-content/uploads/2020/03/MC-TD.pdf>




Dynamic programming

Monte Carlo



Temporal difference





# Summary

---

- Temporal-difference methods update values after each transition
- Uses samples like MC, but bootstraps like DP
- Prediction: One-step ( $TD(0)$ ) for online updates
- Batch updating can be done offline for fixed set of samples
- Control: On-policy (SARSA) vs off-policy (Q-learning)
- SARSA learns values of behavior policy, utilities reflect exploration
- Q-learning learns values of target policy, producing optimal policy